

DAT560 Generative AI

Recurrent Neural Networks

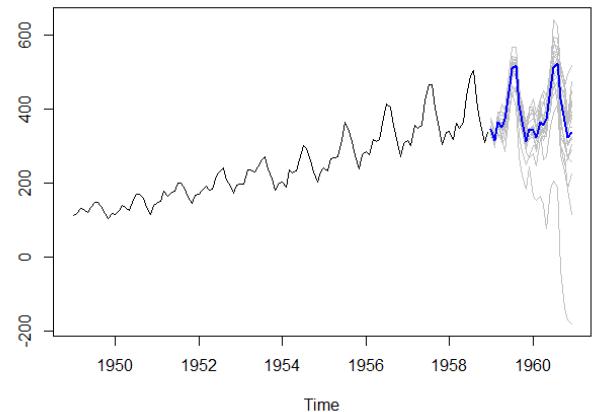


Modelling sequential Data

- What is sequential data ?
 - Data where the order matters – dependencies
- Language Modelling
 - Predict the next word
 - Which word comes next ? “the boy **crossed** the _____”
 - Highly likely : *road, street, highway*
 - Less likely : *pizza, sugar,..*
 - If you understand language well you can generate sensible statements that are grammatically correct and capture world knowledge
- Time Series forecasting

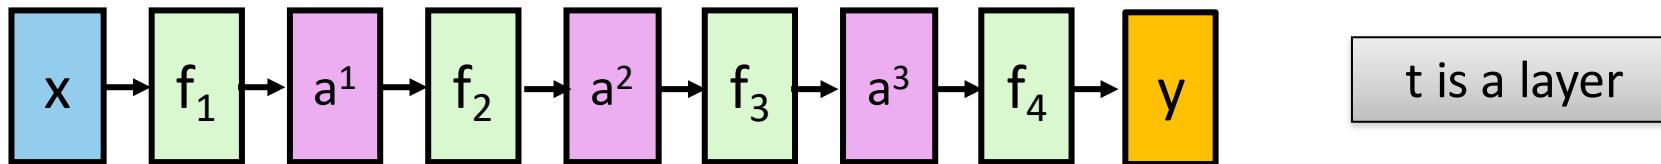
The cat crossed the street, while eating cheese, that was _____

The cat _____ the street that was flat

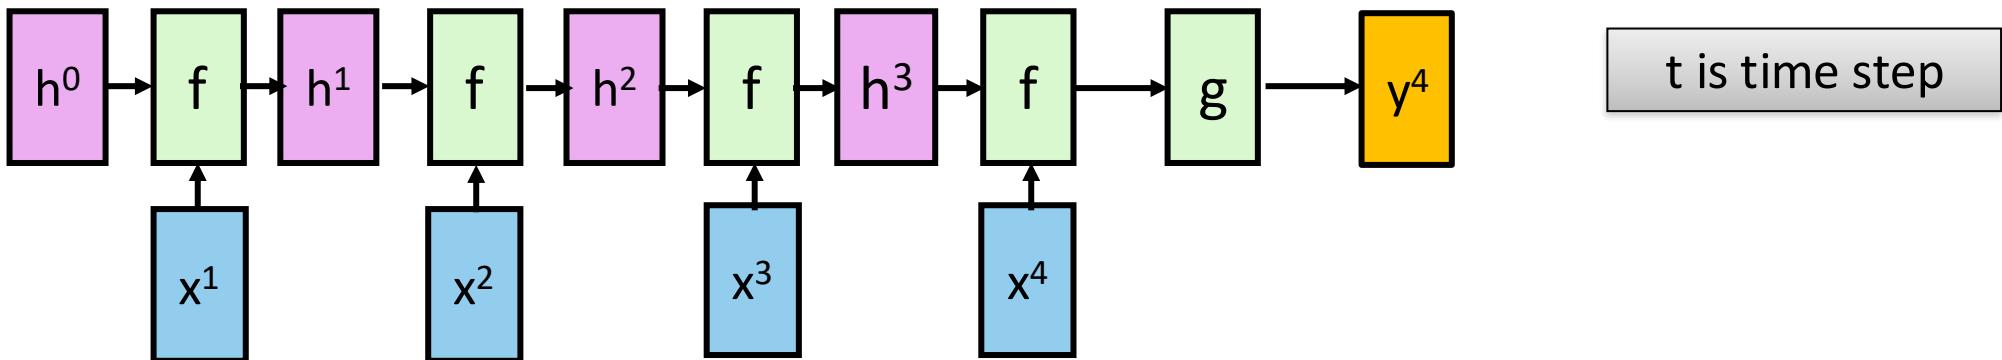


Feedforward vs Recurrent Nets

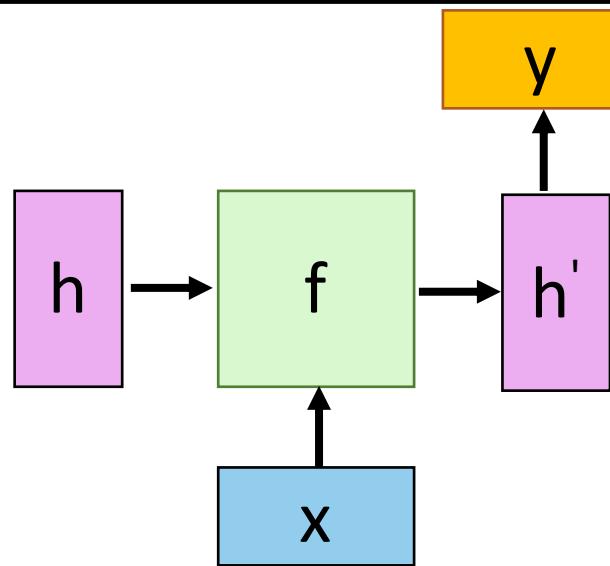
1. Feedforward network does not have input at each step
2. Feedforward network has different parameters for each layer



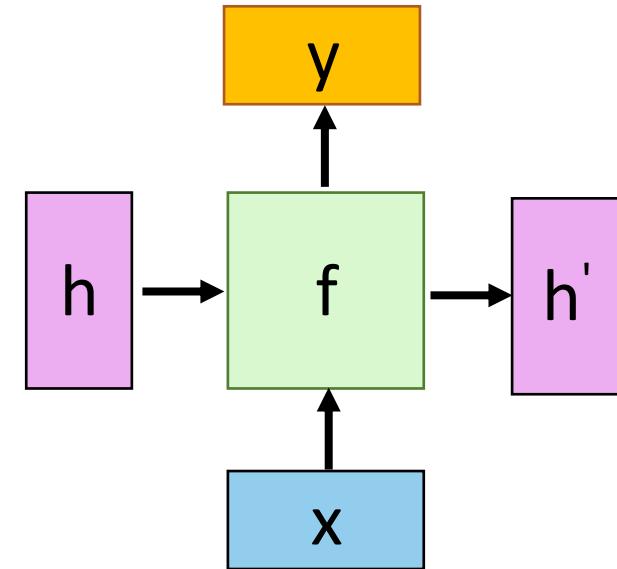
$$a^t = f_t(a^{t-1}) = \sigma(W^t a^{t-1} + b^t)$$



RNN Unit Computation

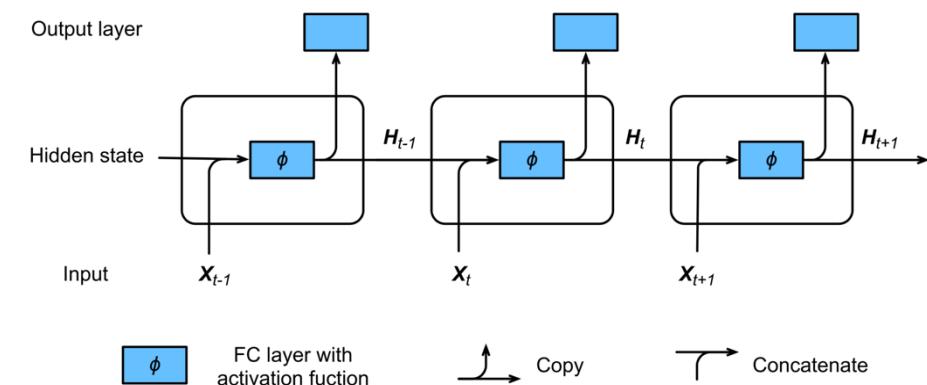


Note, y is computed from h'



$$h' = \sigma_{\text{sigmoid}}(W h + U x)$$

$$y = \sigma_{\text{softmax}}(W^o h')$$

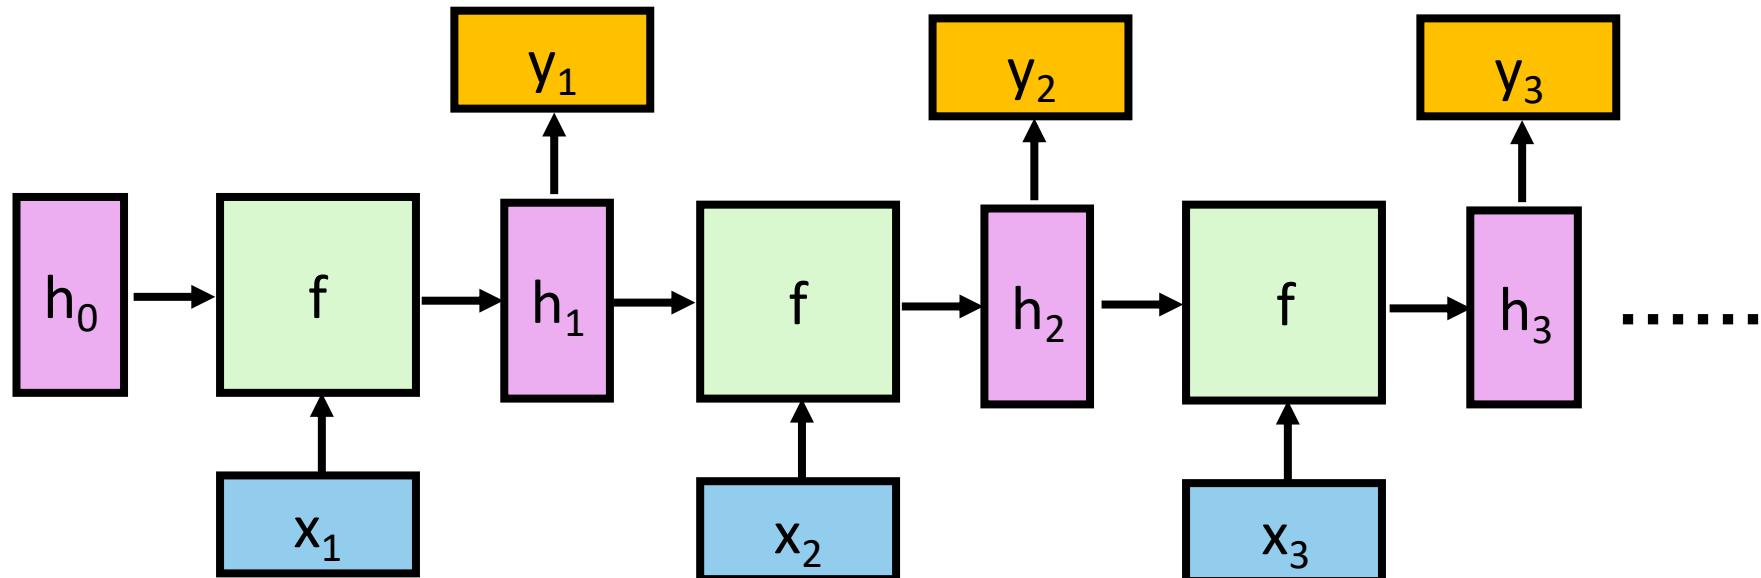


RNN Cell Diagram

Recurrent Neural Network

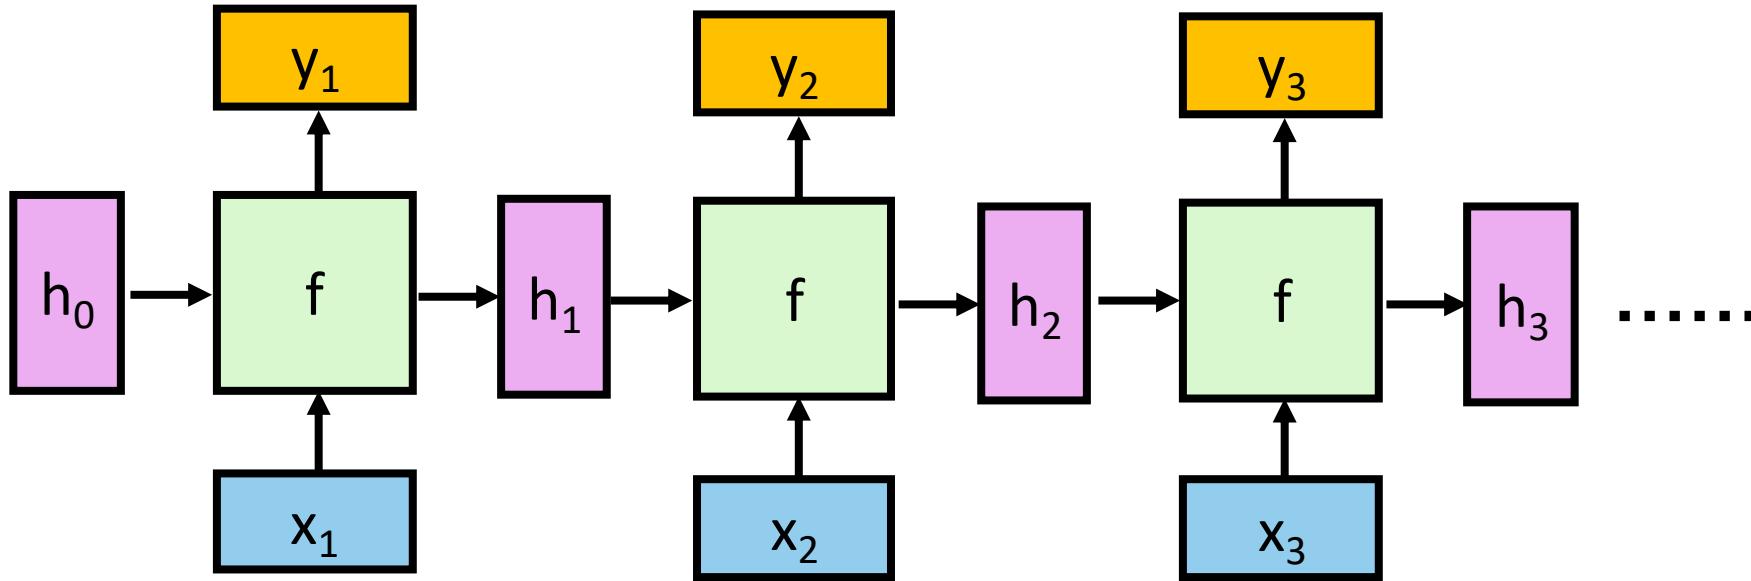
- Given function f : $h', y = f(h, x)$

h and h' are vectors with the same dimension



No matter how long the input/output sequence is, we only need one function f . If f 's are different, then it becomes a feedforward NN. This may be treated as another compression from fully connected network.

Forward Pass



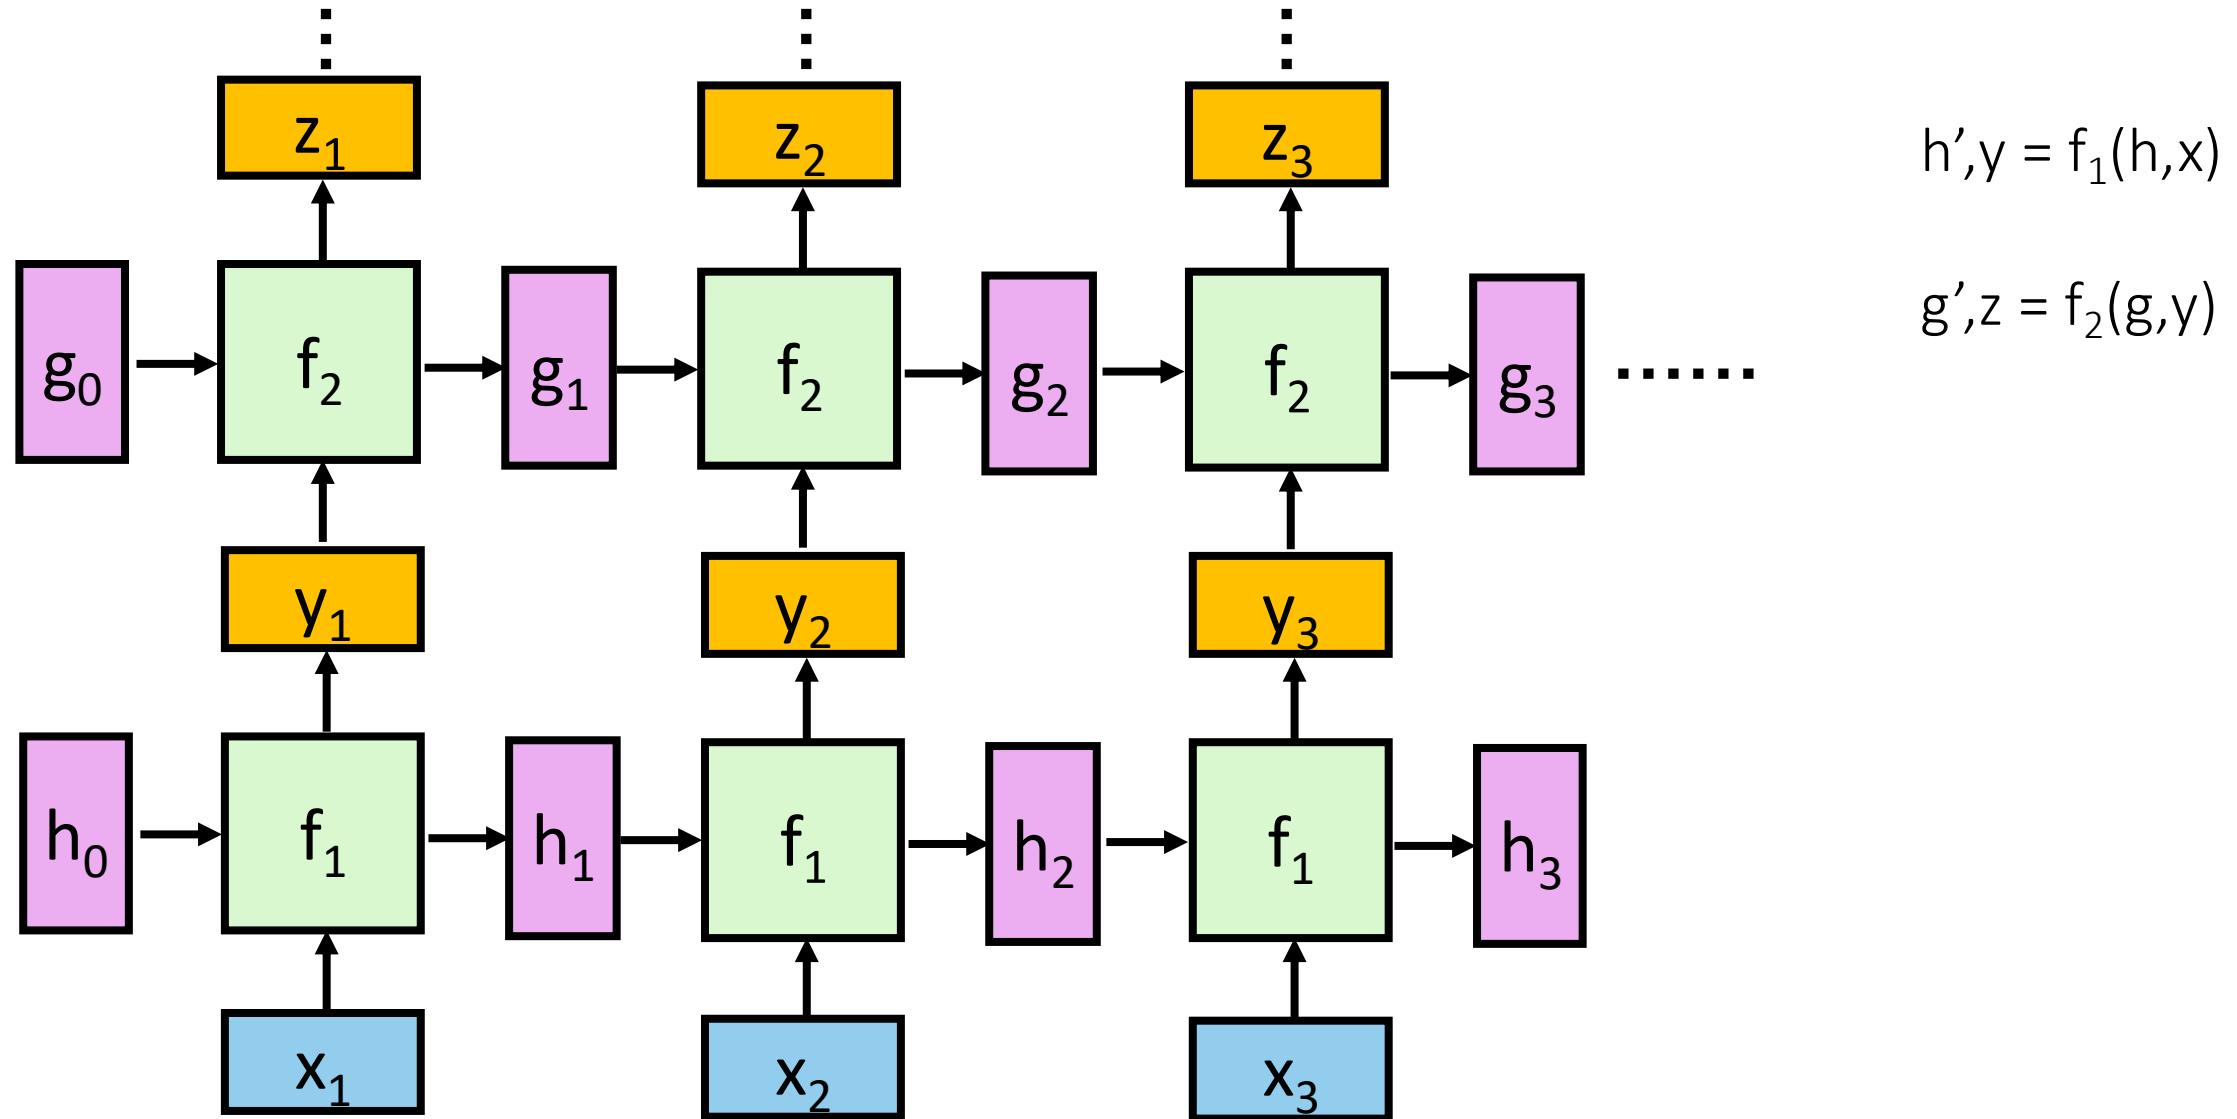
$$h_1 = f(W^T h_0 + U^T x_1)$$

$$h_2 = f(W^T f(W^T h_0 + U^T x_1) + U^T x_2)$$

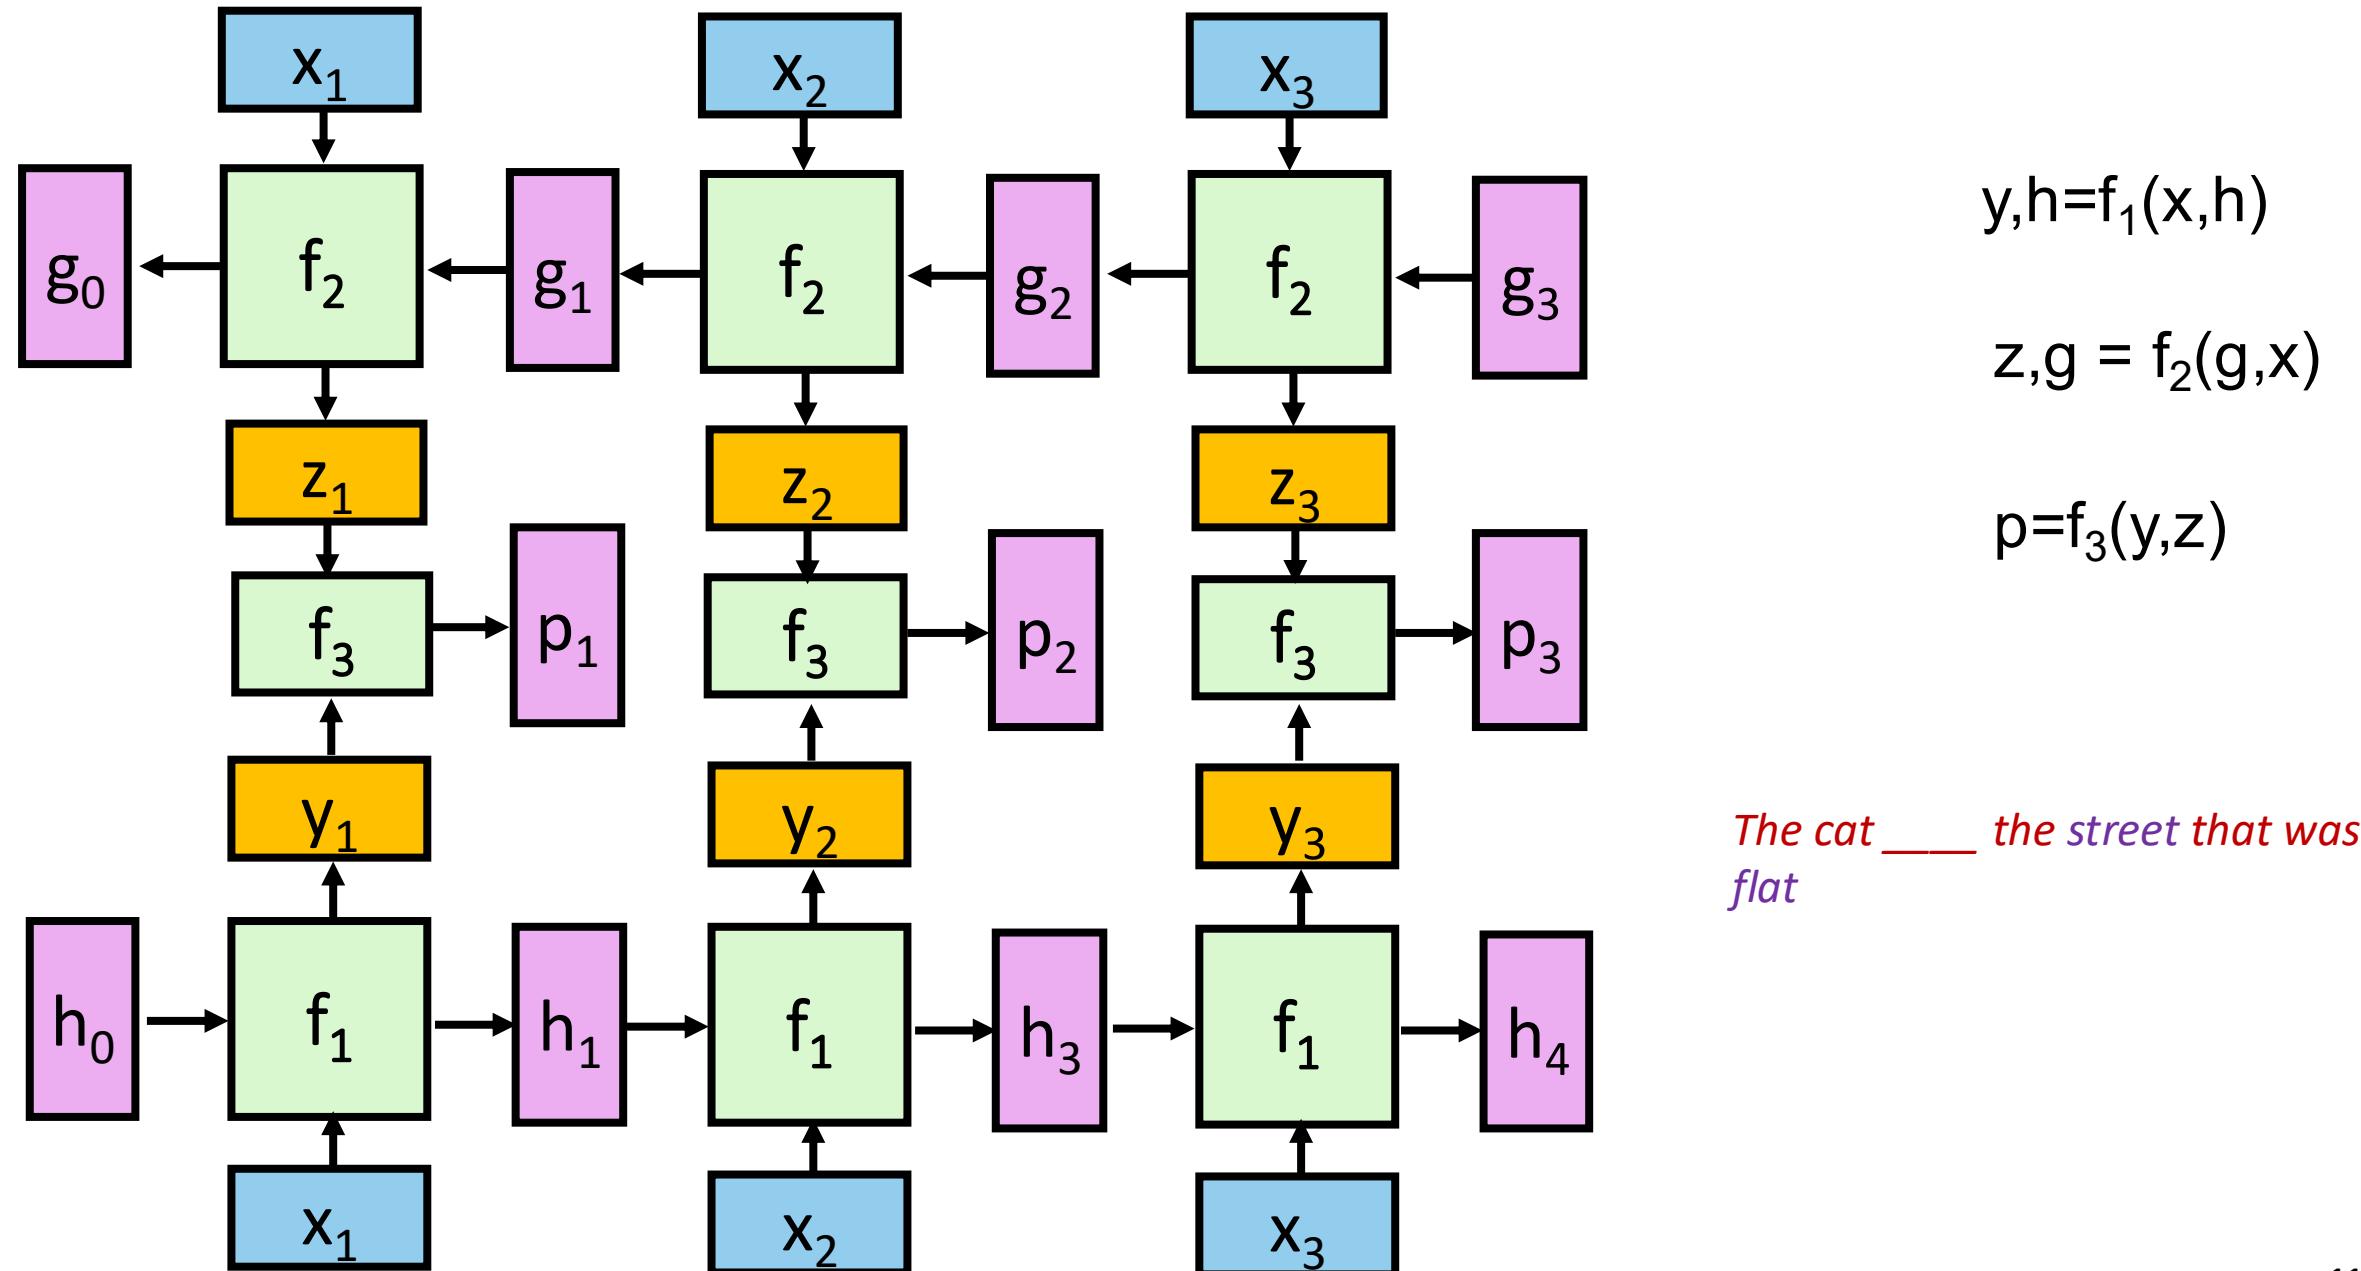
$$h_3 = f(W^T f(W^T f(W^T h_0 + U^T x_1) + U^T x_2) + U^T x_3)$$

Weights are shared

Deep RNN

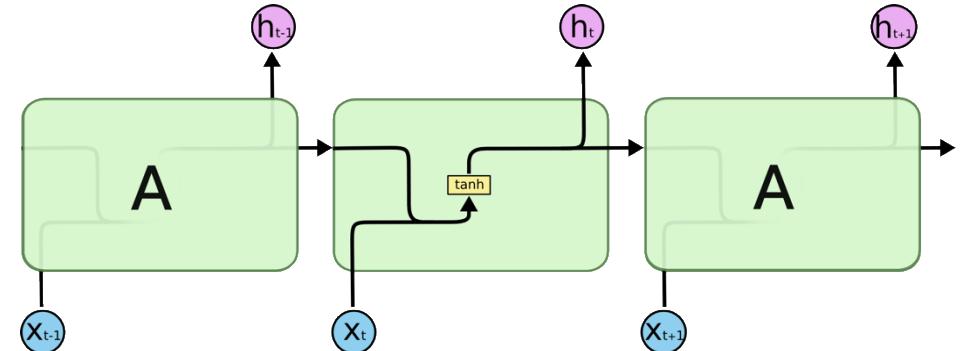


Bi- Directional RNN

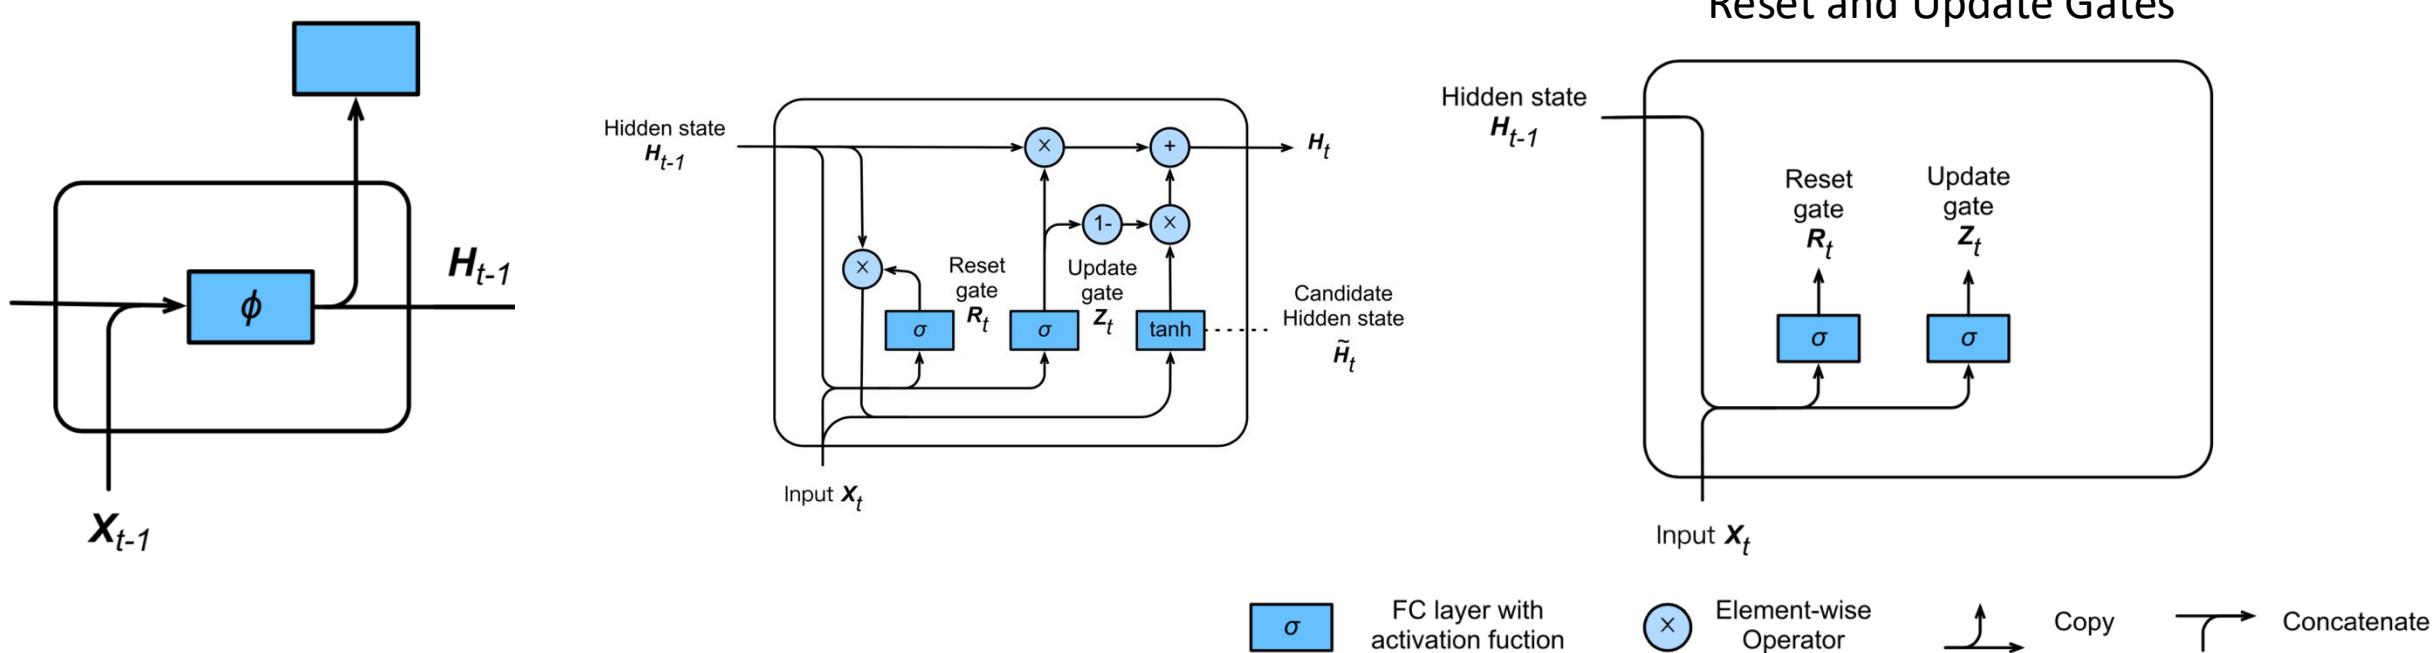


Problems with vanilla RNNs

- Inability to capture long-term dependencies
 - When dealing with a time series, it tends to forget old information. When there is a distant relationship of unknown length, we wish to have a “memory” to it.
- Vanishing and Exploding gradients
 - Weights either become zero or explode due to products of partial differentials
- Slow inference



Gated Recurrent Unit



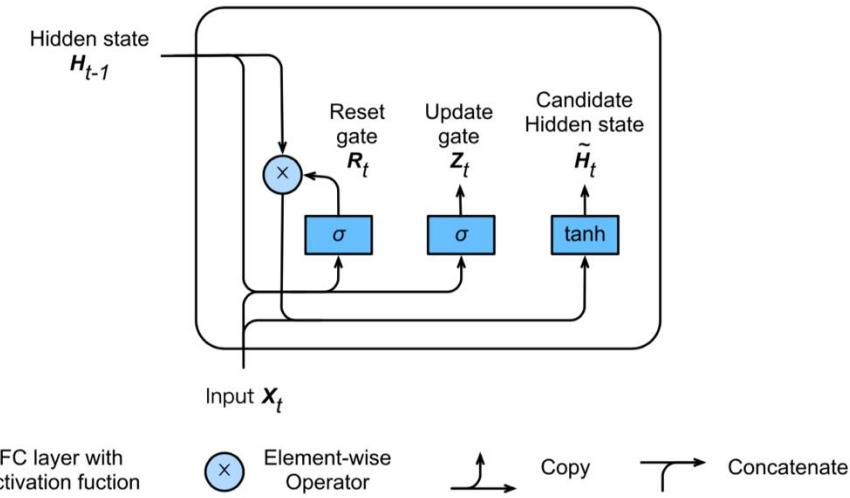
GRUs have the following two distinguishing features:

- **Reset gates** help capture short-term dependencies in time series.
- **Update gates** help capture long-term dependencies in time series.

$$\begin{aligned} \mathbf{R}_t &= \sigma(\mathbf{X}_t \mathbf{W}_{xr} + \mathbf{H}_{t-1} \mathbf{W}_{hr} + \mathbf{b}_r) \\ \mathbf{Z}_t &= \sigma(\mathbf{X}_t \mathbf{W}_{xz} + \mathbf{H}_{t-1} \mathbf{W}_{hz} + \mathbf{b}_z) \end{aligned}$$

Reset gate

- If we want to be able to reduce the influence of previous states
 - multiply H_{t-1} with R_t elementwise
- Whenever the entries in R_t are close to 1 we recover a conventional deep RNN.
- For all entries of R_t that are close to 0 the hidden state is the result of an MLP with X_t as input
- Any pre-existing hidden state is thus ‘reset’ to defaults. This leads to the following candidate for a new hidden state (it is a candidate since we still need to incorporate the action of the update gate).



$$\mathbf{H}_t = \tanh(\mathbf{X}_t \mathbf{W}_{xh} + \mathbf{H}_{t-1} \mathbf{W}_{hh} + \mathbf{b}_h)$$

$$\tilde{\mathbf{H}}_t = \tanh(\mathbf{X}_t \mathbf{W}_{xh} + (\mathbf{R}_t \odot \mathbf{H}_{t-1}) \mathbf{W}_{hh} + \mathbf{b}_h)$$

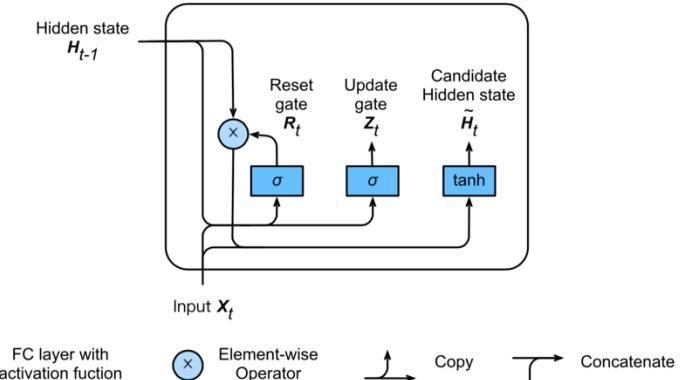
Nonlinearity (Tanh) to ensure
that the values of the hidden
state (-1, 1)

Update gate

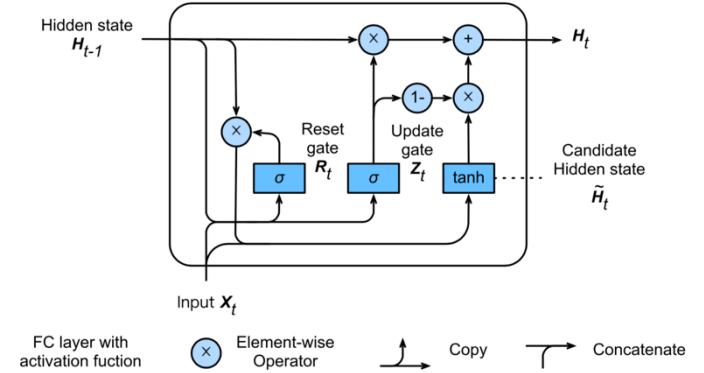
- Determines the extent to which the new state H_t is just the old state and by how much the new candidate state is used.

$$H_t = Z_t \odot H_{t-1} + (1 - Z_t) \odot \tilde{H}_t$$

- Whenever the update gate is close to 1
 - we simply retain the old state.
 - In this case the information from X_t is essentially ignored, effectively skipping time step t in the dependency chain
- Whenever it is close to 0
 - the new latent state H_t approaches the candidate latent state \tilde{H}_t .
 - Helps cope with the vanishing gradient problem in RNNs and better capture dependencies for time series with large time step distances.



$$H_t = Z_t \odot H_{t-1} + (1 - Z_t) \odot \tilde{H}_t$$



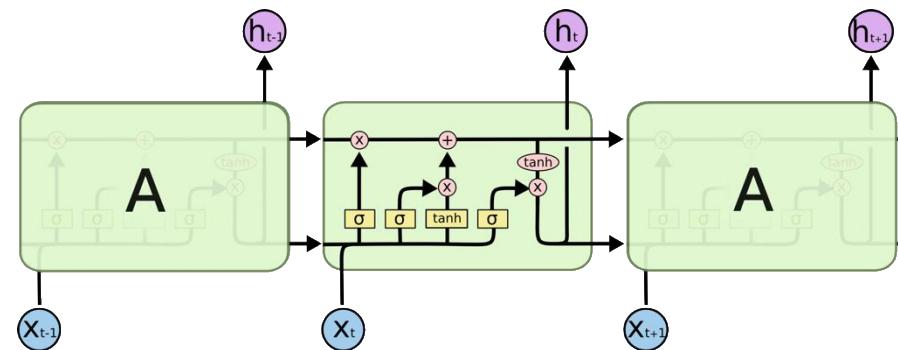
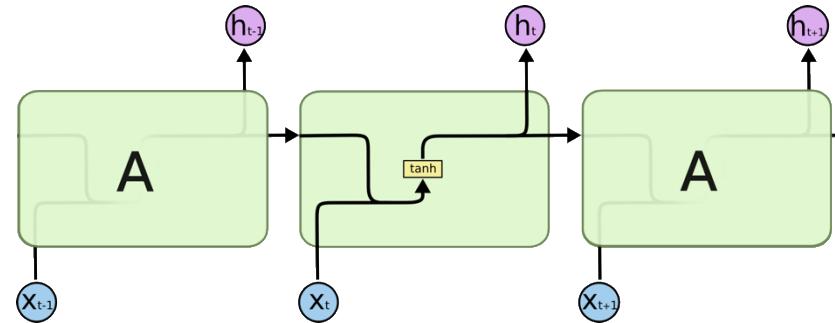
LSTM Networks

- Long Short Term Memory networks – usually just called “LSTMs”
- LSTM was first proposed in 1997 by Sepp Hochreiter and Jürgen Schmidhuber (Neural Computation. 9 (8): 1735–1780.)
 - Submission to NIPS was rejected in 1997!
- It was used by (until Transformers came in)
 - Microsoft for conversational speech recognition
 - Google for machine translation
 - Apple for siri (on iphones)
 - IBM, Baidu, Samsung and so on...



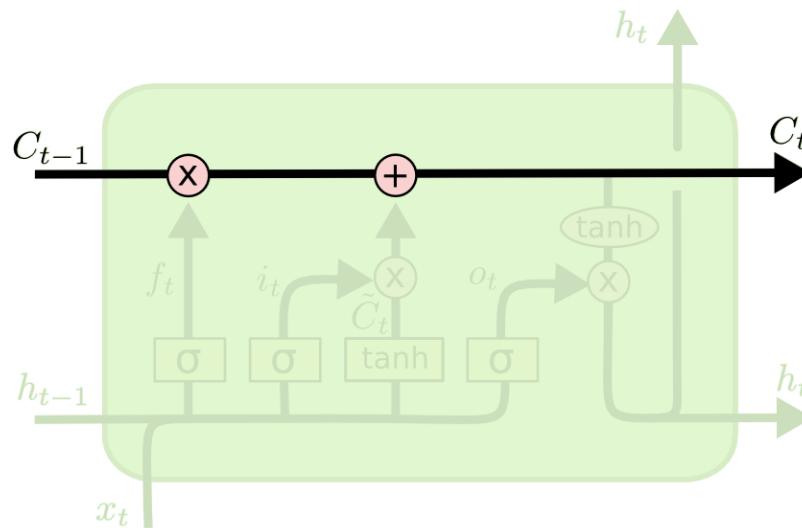
LSTM Networks

- LSTMs are designed to overcome vanishing gradient problems of RNN



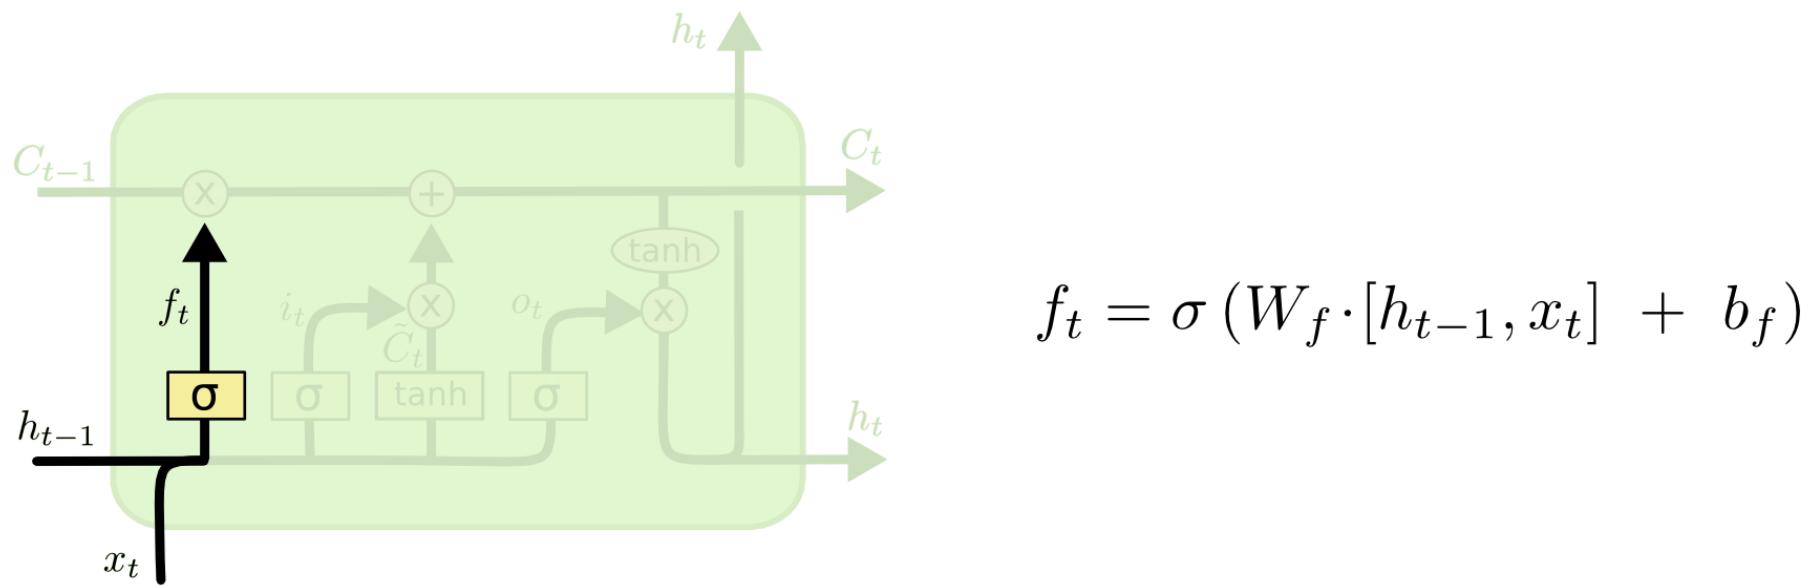
The Core Idea Behind LSTMs

- The key to LSTMs is the cell state, the horizontal line running through the top of the diagram.
- The cell state is kind of like a conveyor belt. It runs straight down the entire chain, with only some minor linear interactions. It's very easy for information to just flow along it unchanged.



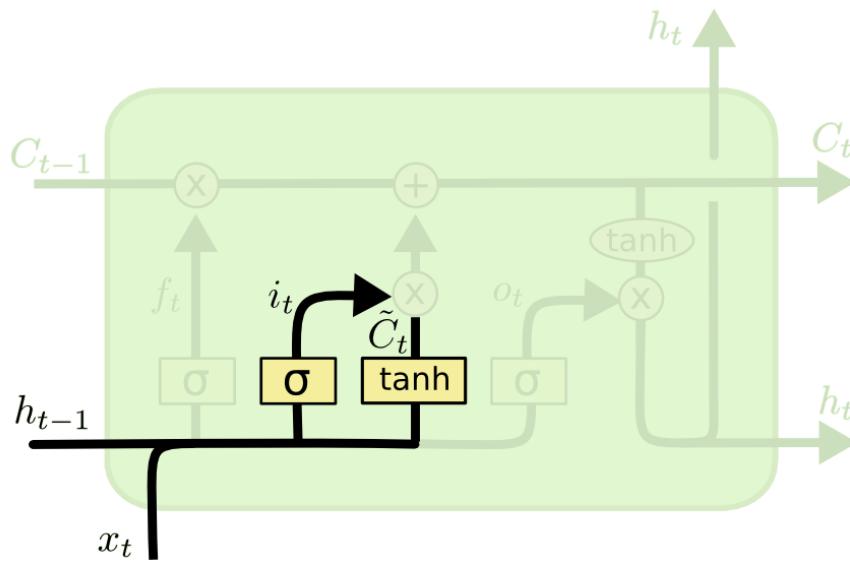
Step-by-Step LSTM Walk Through

- Forget gate: Decides what information to throw away from the cell state
- f_t is between 0 and 1 due to sigmoid
- 1 represents “completely keep this” while a 0 represents “completely get rid of this.”



Step-by-Step LSTM Walk Through

- Input gate: Decides what new information to store in the cell state.
- This has two parts.
 - a sigmoid layer called the “input gate layer”
 - a tanh layer creates a vector of new candidate values, \tilde{C}_t , that could be added to the state

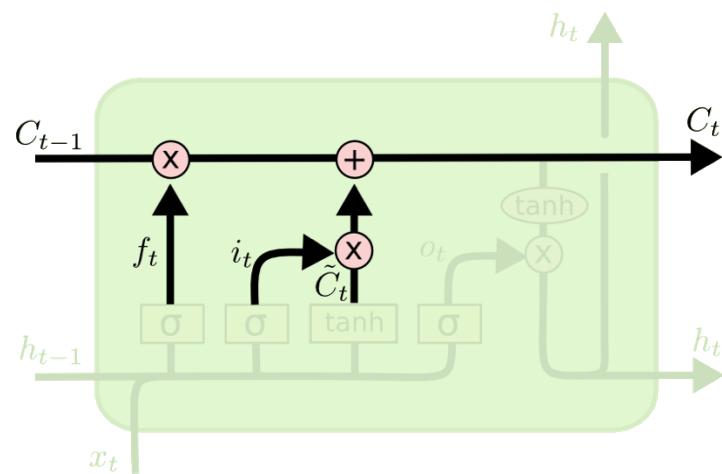


$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Step-by-Step LSTM Walk Through

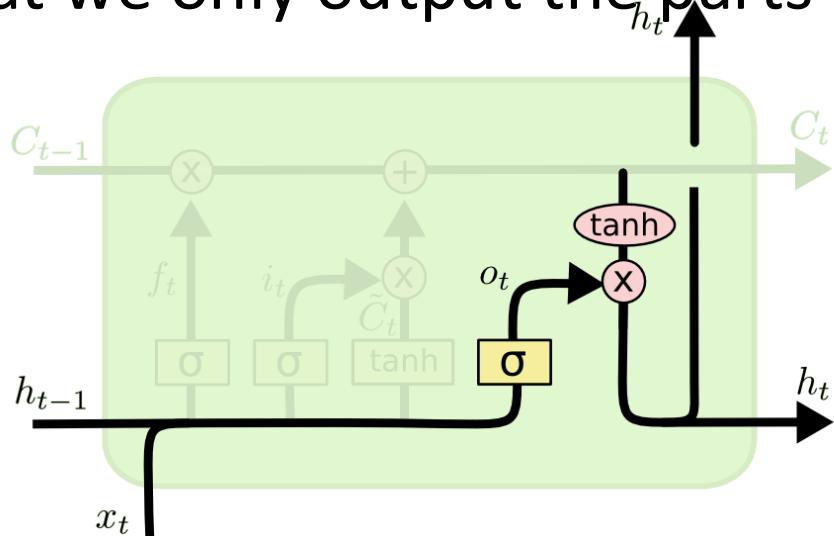
- Next step is to update the old cell state, C_{t-1} , into the new cell state C_t
- We multiply the old state by forget gate output f_t , Then we add $i_t * \tilde{C}_t$.



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Step-by-Step LSTM Walk Through

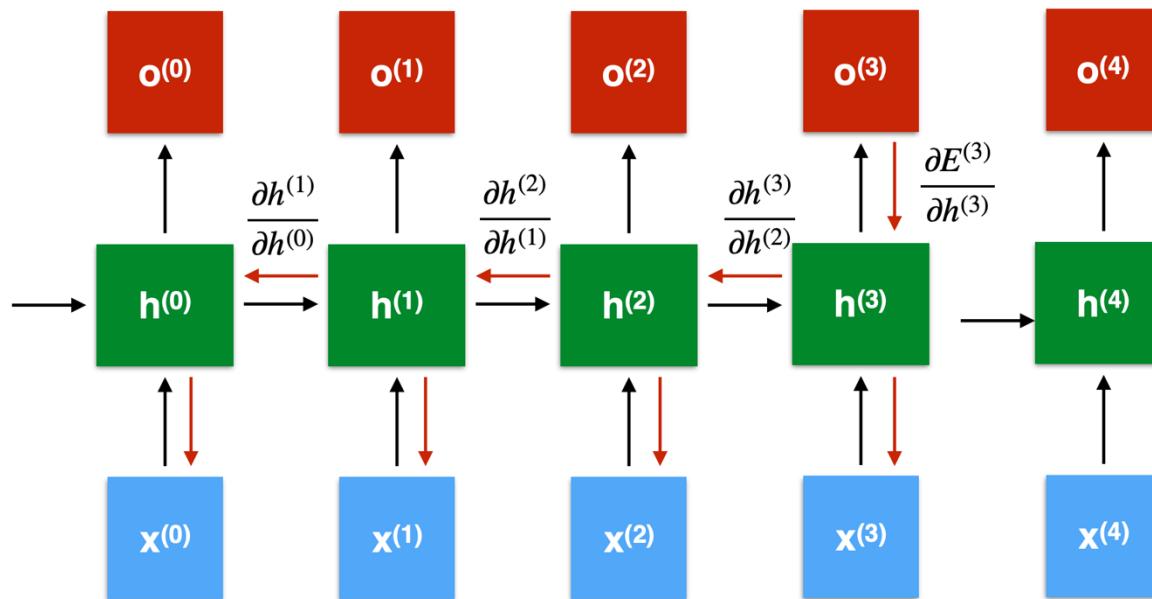
- Finally the output will be based on our cell state, but will be a filtered version.
- First, we run a sigmoid layer which decides what parts of the cell state we're going to output.
- Then, we put the cell state through tanh (to push the values to be between -1 and 1) and multiply it by the output of the sigmoid gate, so that we only output the parts we decided to.



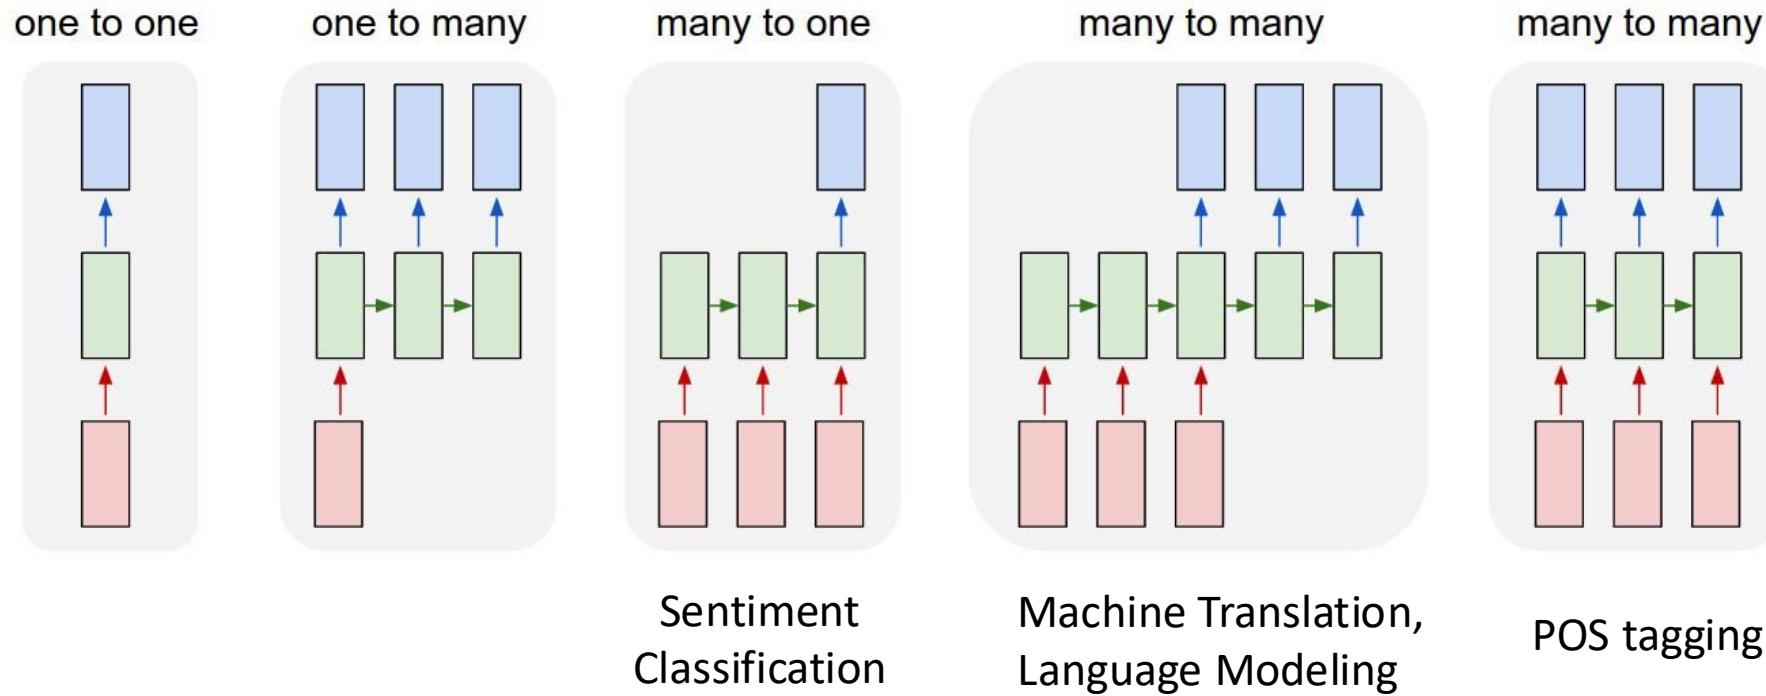
$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t * \tanh (C_t)$$

Back Propagation through Time

- One of the methods used to train RNNs
- The unfolded network (used during forward pass) is treated as one big feed-forward network
- This unfolded network accepts the whole time series as input
- The weight updates are computed for each copy in the unfolded network, then summed (or averaged) and then applied to the RNN weights



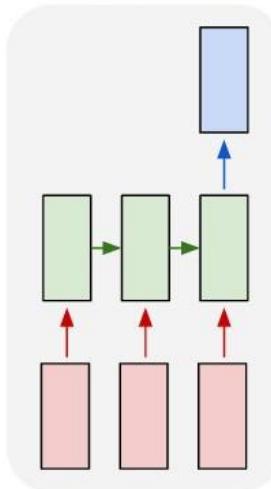
Different Scenarios



Sentiment Classification

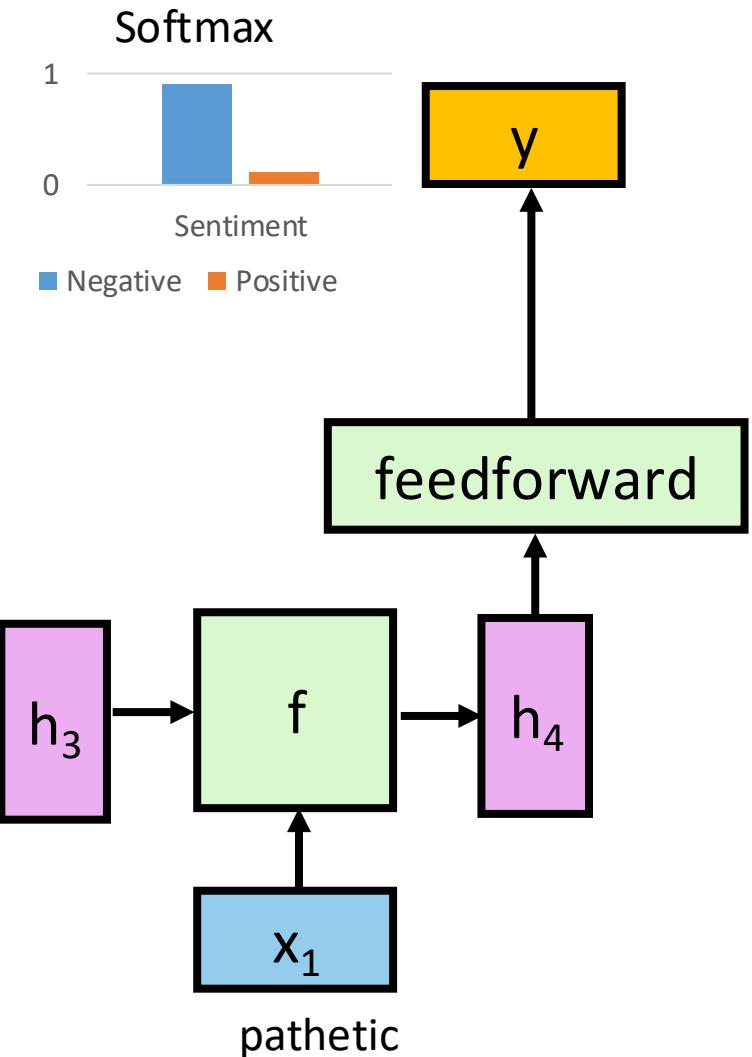
- **Task:** Given a review (natural language text) classify it as a positive or a negative sentiment
 - “the movie is pathetic” → {positive, negative}
- Input: pre-trained word embeddings
- Each cell is a GRU cell
- Loss function: Cross entropy loss

many to one



Simple RNN for Sentiment Classification

- Task: Given a review (natural language text) classify it as a positive or a negative sentiment
 - “the movie is pathetic” → {positive, negative}
- Input: pre-trained word embeddings
- Each cell is a GRU cell
- Loss function: Cross entropy loss



References

- Luis Serrano, A Friendly Introduction to Recurrent Neural Networks, <https://www.youtube.com/watch?v=UNmqTiOnRfg>, Aug. 2018
- Brandon Rohrer, Recurrent Neural Networks (RNN) and Long, Short-Term Memory (LSTM), <https://www.youtube.com/watch?v=WCUNPb-5EYI>, Jun. 2017
- Denny Britz, Recurrent Neural Networks Tutorial, <http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>, Sept. 2015 (Implementation)
- Colah's blog, Understanding LSTM Networks, <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>, Aug. 2015
- <https://distill.pub/2019/memorization-in-rnns/>
- <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Slides credit

- Avishek Anand

